

Atomic Operations

Dmitri Nesteruk

@dnesteruk



Overview

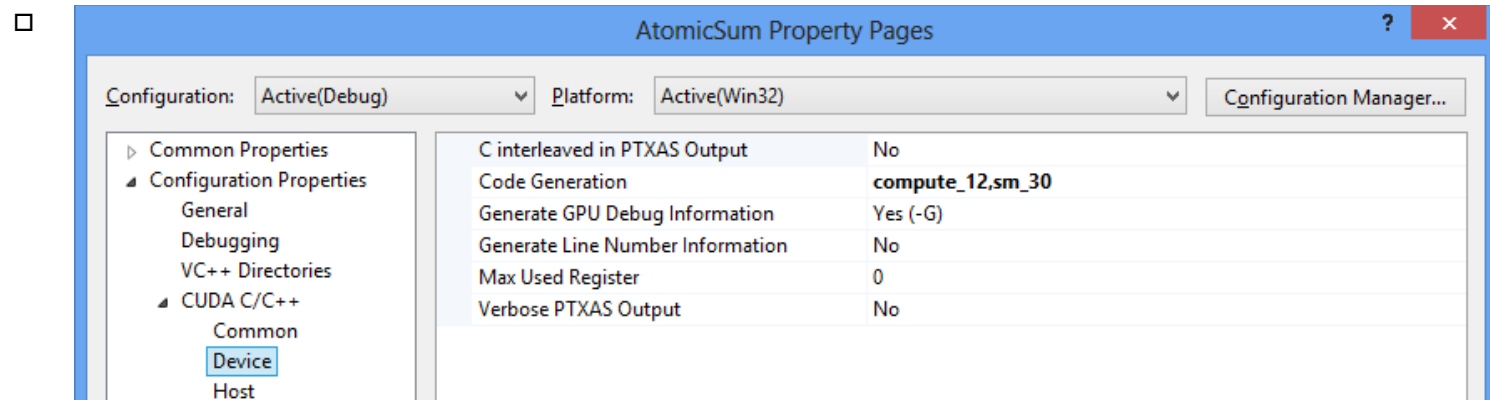
- Why atomics?
- Atomic Functions
- Atomic Sum
- Monte Carlo Pi

Why Atomics?

- **x++ is a read-modify-write operation**
 - Read x into a register
 - Increment register value
 - Write register back into x
 - Effectively { temp = x; temp = temp+1; x = temp; }
- **If two threads do x++**
 - Each thread has its own temp (say t1 and t2)
 - { t1 = x; t1 = t1+1; x = t1; }
 { t2 = x; t2 = t2+1; x = t2; }
 - Race condition: the thread that writes to x first wins
 - Whoever wins, x gets incremented only *once*


Atomic Functions

- Problem: many threads accessing the same memory location
- Atomic operations ensure that only one thread can access the location
- Grid scope!
- `atomicOP(x,y)`
 - `t1 = *x; // read`
 - `t2 = t1 OP y; // modify`
 - `*a = t2; // write`
- Atomics need to be configured
 - `#include "sm_20_atomic_functions.h"`



Monte Carlo Pi

- Evaluate the value of π numerically
- Area of circle $S_c = \pi r^2$
- Area of square $S_s = 2r \times 2r = 4r^2$
- $\therefore \pi = \frac{4S_c}{S_s}$
- Generate random points within the square
- Count the number of points in the circle
 - $\sqrt{x^2 + y^2} < r$
- $\pi \approx \frac{4 \times \text{points in circle}}{\text{total number of points}}$



Summary

- **Atomic operations ensure operations on a variable cannot be interrupted by a different thread**
- **CUDA supports several atomic operations**
 - `atomicAdd()`
 - `atomicOr()`
 - `atomicMin()`
 - ... and others
- **Atomics incur a heavy performance penalty**